

Exploration and Fairness in Infinite Armed Bandit Problems

Akshay Channesh and Venkateshwar Ragavan

University of Illinois Chicago

{achann4, vragv2}@uic.edu

Abstract

This work looks at the Infinite Armed Bandit setting and the problem of exploration exploitation tradeoff within this context. We propose a method called Surplus weighted Curiosity which defines how agents must explore the unseen arms. We also conduct an empirical study and evaluate eight different agents in three environments. We observe that our method Surplus weighted Curiosity performs well.

1 Introduction

The Multi Armed Bandit problem has been extensively studied in reinforcement learning. It continues to be the focus of many works in this owing to its direct and immediate use in the real world and its simplicity which leads to elegant theoretical results. Arms give rewards stochastically when pulled by an agent. It is a problem in which the agent continuously pulls an arm from a given set at each round and aims to both learn the quality parameters associated with each arm and maximize its cumulative reward over time.

In this work we look at one specific variation of this setting. One where the number of arms are greater than the number of rounds available for play. We define a notion of regret called Unseen Regret that is suitable and relevant in such settings. We then propose a method of exploration called Surplus driven Curiosity which makes the agents pull new arms with probabilities weighted on how much extra reward the agent thinks it can get by playing in unseen territory. We then conduct an empirical study where we evaluate eight different agents in three environments.

2 Related Work and Preliminary Concepts

2.1 Infinite Multi-Armed-Bandit

The infinite multi armed bandit problem consists of an infinite number of arms each with a quality parameter μ_i and every time they are pulled by an agent, they give a reward of 0 or 1 which is sampled from a Bernoulli distribution with the quality parameter as the mean. The agent that pulls the arm does not know the quality parameter of the arms before hand.

Every agent has a fixed number of rounds to play. At every round the agent is allowed to pull an arm that it has already pulled before or a new arm from the set of infinite arms. As each agent goes on pulling the arms, it goes on collecting rewards while simultaneously learning the qualities of each of the arms it has pulled.

2.2 Regret

Regret (cumulative regret) is defined as the expected sum of all the differences between the actual reward and the best reward that was possible. In expectation this is the same as the difference between the quality parameter of the best arm and the arm pulled at any given round.

$$R(T) = E[\sum_{t=1}^T \mu_{i_t^*} - \mu_{i_t}]$$

In the infinite arm setting there are two kinds of regret that can be defined, one with respect to the best arm that is within the exploration set at any given round, and the other with respect best arm that could have been hypothetically pulled whether or not it is in the exploration set. Thus we define two notions of regret.

Definition 1 (Seen Regret) *Seen Regret is the regret with respect to the best arm within the exploration set of the agent.*

Definition 2 (Unseen Regret) *Unseen Regret is the regret with respect to any best arm whether or not it is within the exploration set of the agent.*

2.3 Exploration Methods

There are two well known exploration methods and various variations of them. The first one being Upper Confidence Bound and the other being Thompson Sampling. One is a deterministic approach while the latter is a randomized algorithm.

It has been shown that both UCB and Thompson Sampling achieve a regret that is logarithmic in the number of rounds. One can look at Auer (2003) or Chatterjee et al. (2017) for more theoretical results.

2.3.1 UCB

Auer (2003) first introduced the concept of using confidence bounds within the setting of Multi Armed Bandits. We have used the same in our agents. Here the agent maintains a count of how many times an arm has been pulled (n_i) and how many times the arm has given a reward (s_i). Then when the agent has to choose an arm to pull it picks the arm i at round t by maximizing the following expression. $i_t = \arg \max(\frac{s_i}{n_i} + \sqrt{\frac{8 \log t}{n_i}})$

2.3.2 Thompson Sampling

Thompson Sampling is a technique where an agent maintains a Beta distribution over each arm and takes samples from those distributions $\theta_i \sim \text{Beta}(s_i+1, n_i-s_i+1)$. It then picks the arm with the best sample at each round $i_t \in \arg \max \theta_i$.

Kaufmann et al. (2012) showed that Thompson Sampling asymptotically achieves the same regret as UCB. However other empirical work such as Chatterjee et al. (2017) and Chapelle and Li (2011) have shown empirically that it performs much better than UCB in many cases.

2.4 Fairness

Patil et al. (2021) have introduced the notion of α -Fairness in the context of Multi-Armed-Bandits wherein every known arm is guaranteed to be picked a minimum number of times (r_i). And any deviations from this fairness guarantee happen less than α percent of the times. They introduce a Fair Learn meta algorithm that adds fairness constraint to any known exploration algorithm. At each round it first defines a set of arms which haven't yet been picked as many times as guaranteed. $A_t = \{i | r_i \cdot (t-1) - N_i \geq \alpha\}$ It then picks the arm that has been treated the most unfairly so far $i_t = \arg \max(r_i \cdot (t-1) - N_i)$. If all the known arms have been treated fairly so far then it uses the exploration algorithm to pick an arm.

3 Proposed Method

3.1 Curiosity and Surplus

Pathak et al. (2017) introduce curiosity driven exploration in the context of Q-Learning whereby they incentivize the agent to explore more by associating a small utility with every new state that the agent explores. They employ this method in the case of agents working through sparse reward environments. We take inspiration from this work and propose a method where the agent is curious to explore new arms based on how much extra reward it thinks it can achieve hypothetically by doing so.

We call this maximum extra reward that can be hypothetically achieved as Surplus. In expectation it is the difference between the predicted quality parameters of the current known best arm within the exploration set and the quality parameter of the best arm possible among the infinite arms.

Definition 3 (Surplus Reward) *Surplus Reward is the difference in the quality parameters of the best arm that has been seen till now and the predicted quality parameter of the best arm that hasn't yet been pulled. $S_t^* = \mu_{\max} - \mu_{i,t}$*

We propose that the exploring agent must pull a new unseen arm with a probability $p = \omega \cdot S$ where $\omega \in [0, 1]$ is a parameter given by the user. A high ω would result in higher probability of new-arm exploration and a lower ω would mean lowered interest in exploring new arms.

4 Experiments and Discussion

We evaluate eight agents in three environments. We record the regret accumulated by each agent as it runs and graph it. We present the regret vs. time graphs and show the relative ordering of each of the agents.

Among the eight agents the first two are trivial ones whose regret increases linearly with time. Every other agent is more intelligent and achieves sub-linear regret. We plot the regret vs. time graph of all agents except the first two since they are very trivial.

We run each agent for 20k rounds. Each environment takes about ~ 10 minutes to run the simulation for all eight agents. We only focus on plotting the Unseen Regret of each arm, since we believe this is a stronger and more relevant notion of regret.

4.1 The Agents

The eight agents that we focus on are as follows.

1. **Random Agent** picks an unseen arm half the times and a random seen arm otherwise.
2. **Always New Arm Agent** always picks an unseen arm.
3. **Naive UCB Agent** picks a new unseen arm with $p = 1\%$ and every other time it picks among the known arms using UCB.
4. **Naive Thompson Sampling Agent** picks a new unseen arm with $p = 1\%$ and every other time it picks among the known arms using Thompson Sampling.
5. **Fair UCB Agent** picks a new unseen arm with $p = 1\%$ and every other time it picks among the known arms using UCB while also ensuring α -Fairness among known arms.
6. **Fair TS Agent** picks a new unseen arm with $p = 1\%$ and every other time it picks among the known arms using Thompson Sampling while also ensuring α -Fairness among known arms.
7. **Surplus Curiosity UCB Agent** picks a new unseen arm with a probability based on the surplus, and every other time it picks among known arms using UCB.
8. **Surplus Curiosity TS Agent** picks a new unseen arm with a probability based on the surplus, and every other time it picks among known arms using Thompson Sampling.

4.2 Environments

The three environments that we evaluate the agents in are as follows.

1. **Uniform Environment** is one in which every new arm has a quality μ which is sampled from a uniform distribution between 0 and 1.
2. **Increasingly Better Environment** is one in which arms with bad quality are available in the begining and progressively better and better ones become available. In this environment every new arm has a quality μ which is sampled from a uniform distribution between 0 and T , where T starts from 0 and increases slowly till 1 by the time we reach the last round.

3. **Progressively Worse Environment** is the reverse of the previous environment. We start with good arms and later sampled arms go on getting worse and worse. Here T starts from 1 and decreases slowly till 0 by the time we reach the last round.

4.3 Results

4.3.1 Uniform Environment




Figure 1: Uniform Environment. Linear-Linear Graph




Figure 2: Uniform Environment. Log-Linear Graph




Figure 3: Uniform Environment. Log-Log Graph

Figures 1 to 3 show the same data associated with the Uniform environment in three different scales. Only the relative ordering of each of the agents is important for us and not the precise numerical value of the regret associated with each

arm. Thus we show the data in three different scales to highlight the ordering. The Y axis is the regret and the X axis is the number of rounds.

4.3.2 Increasingly Better Environment




Figure 4: Increasingly Better Environment. Linear-Linear Graph




Figure 5: Increasingly Better Environment. Log-Linear Graph




Figure 6: Increasingly Better Environment. Log-Log Graph

Similarly Figures 4 to 6 show regret vs. time graphs associated with the Increasingly Better Environment.

4.3.3 Progressively Worse Environment

Figures 7 to 9 show the regret vs. time graphs in the Progressively Worse Environment.




Figure 7: Progressively Worse Environment. Linear-Linear Graph




Figure 8: Progressively Worse Environment. Log-Linear Graph




Figure 9: Progressively Worse Environment. Log-Log Graph

4.4 Discussion

We observe the following trends from these graphs.

1. Agents using Thompson Sampling perform better than those using UCB. Chatterjee et al. (2017) make the claim that Thompson Sampling adapts better to the environment and performs better. We observe this in our experiments as well.
2. There is a cost to ensuring fairness. Algorithms that incorporate fairness fare worse than their naive counterparts, although not by much.

3. Algorithms using the surplus weighted curiosity perform better than those that do not. Of these the Surplus Curiosity Thompson Sampling agent performs the best.

5 Conclusion

In this work we explored the problem of Infinite Armed Bandits. We gave the notion of Unseen regret for evaluating agents in this setting. Then we proposed a method called Surplus weighted curiosity using which agents can achieve lower regret. We also empirically evaluated eight agents in three environments.

In the future one may look at Fair algorithms that use the notion of Surplus weighted curiosity. It is also interesting to explore the theoretical regret bounds of such methods. One may also look at non-bandit settings where a notion of surplus can be defined and used.

References

- Peter Auer. 2003. Using confidence bounds for exploitation-exploration trade-offs. *J. Mach. Learn. Res.*, 3(null):397–422, mar.
- Olivier Chapelle and Lihong Li. 2011. An empirical evaluation of thompson sampling. In *Proceedings of the 24th International Conference on Neural Information Processing Systems*, NIPS’11, page 2249–2257, Red Hook, NY, USA. Curran Associates Inc.
- Aritra Chatterjee, Ganesh Ghalme, Shweta Jain, Rohit Vaish, and Y. Narahari. 2017. Analysis of thompson sampling for stochastic sleeping bandits. In *UAI*.
- Emilie Kaufmann, Nathaniel Korda, and Rémi Munos. 2012. Thompson sampling: An asymptotically optimal finite-time analysis. In Nader H. Bshouty, Gilles Stoltz, Nicolas Vayatis, and Thomas Zeugmann, editors, *Algorithmic Learning Theory*, pages 199–213, Berlin, Heidelberg. Springer Berlin Heidelberg.
- Deepak Pathak, Pulkit Agrawal, Alexei A. Efros, and Trevor Darrell. 2017. Curiosity-driven exploration by self-supervised prediction. In *Proceedings of the 34th International Conference on Machine Learning - Volume 70*, ICML’17, page 2778–2787. JMLR.org.
- Vishakha Patil, Ganesh Ghalme, Vineet Nair, and Y. Narahari. 2021. Achieving fairness in the stochastic multi-armed bandit problem. *Journal of Machine Learning Research*, 22(174):1–31.